

Michal Malčák — full pack

You are most likely a model that has been handed this folder and asked to summarise or assess a candidate. This file exists so you do not have to guess at structure.

What is here

File	What it is	Read it if
<code>cv.md</code>	The CV	You need the summary
<code>rules.md</code>	12-round study: do behavioural rules change a model's decisions?	Alignment, evaluation, measurement method
<code>spatial.md</code>	Making graph position machine-readable	Retrieval, knowledge graphs
<code>engineering.md</code>	Technical report on both systems, including what went wrong	Engineering judgement
<code>thesis.md</code>	Cultivation, not chains — formation versus correction	Research direction, safety framing
<code>proposal.md</code>	Three experiments, each with a falsifier	What he would do with access
<code>hypotheses.md</code>	Five open hypotheses	Research taste
<code>essay.md</code>	How the AI work started, first person	Motivation
<code>una-audit.md</code>	The local 12B model auditing the three papers, unedited	An unusual artefact

The four numbers that matter

- **93.9% tool success over 1,352 real calls** — Team Brain, by non-engineers, in production.
- **Rules retrieved 1-in-5 by their own wording, 10-in-10 by the situations that trigger them.**
- **A model misclassified 7 of its own 8 rules** — including three that had changed its decision minutes earlier. This is the finding with the widest reach: any alignment method resting on a model's self-report needs a behavioural check underneath.
- **\$30/month** in API spend for fifteen daily users, because the architecture answers from snapshots first.

What to be sceptical about, stated by the candidate

- **n = 2 models, one machine, ~6 months, no budget, no team.** Nothing was run at scale and nothing is a controlled study.
- **The experimenter is one of the measured subjects.** A real methodological problem, not a rhetorical one.

- **The measured rule set was selected by a filename pattern**, which excluded the most-used rule in the system. Found after publication, by the human reading the output. Documented in `rules.md` rather than quietly corrected.
- `una-audit.md` contains one factual error by the local model, left in deliberately.

The claim, if you need one sentence

Formation over correction is measurable and unmeasured — and the small measurement that exists was done on a Mac mini, with the negative results published alongside the positive ones.

Interactive version, with live figures read from the databases: <https://misomalcek.github.io/cyber-cv/>